

Engineering 8894/9875:

File abstractions

<https://memorialu.gitlab.io/Engineering/ECE/Teaching/operating-systems/website/lecture/3/>

Recall

Virtualization

Time-sharing: Each user has the impression of working on a dedicated computer

Virtual memory: Each process has the impression of a dedicated address space

How about files?

Files*

What is a file?

- an array of bytes that *persists*
- can be found via a *path*, e.g.:

```
/home/jon/Teaching/8894/website/config.yaml
```

- **don't** have user-meaningful names

* Reference: [OSTEP Ch 39](#)

3 / 15

Comparison with networks: "people get mad"

They *do* have not-terribly meaningful names... we'll explain a bit more about that soon!

Directories

Folders?

Set of entries ($n \rightarrow i$)

n : user-meaningful name

i : a file's *inode*

e.g.: (`foo.c` \rightarrow 925551)



... um, but what's an inode?

inodes


- **probably** (!) short for "index node"
- integer index that names a location on disk
- location can be the start of a file...
- ... or a directory!

5 / 15

This is the _____ name that a file itself *really* has

So we can put _____

Another directory hierarchy



... but so could `/home/jon/workspace/notes!`

But `/home/jon/TODO-notes` is *not* the name of the file!

Directory summary

Directories map names to inodes

inodes identify files and directories

Paths can be used to look up files...

... but files and paths change independently

- file contents can change without changing the path
- path can change without changing the file

Process file abstractions

But how do *processes* access files?

```
new PrintWriter(new BufferedWriter(new FileWriter(name))).write("Hello!");
```

```
std::cout << "Hello, world!" << std::endl;
```

```
fwrite(stdout, "Hello, world!\n");
```

```
char message[] = "Hello, world!\n";  
int fd = STDOUT_FILENO;  
write(fd, message, sizeof(message));
```


Process file abstractions

Each process has a set of integer *file descriptors**

Can use *system calls* to open, close, read, write, etc.

```
int fd = open("/home/jon/hello.txt", O_RDONLY);  
  
/* ... */  
  
write(fd, some_data_bytes, data_length);  
  
/* ... */
```

* We'll come back to what a file descriptor means soon for now, just think of it as a small integer that identifies a particular file.

File I/O system calls

```
/* Open or close a file: */
int    open(const char *, int, ...);
int    close(int);

/* Sequential reading and writing: */
ssize_t read(int, void *, size_t);
ssize_t write(int, const void *, size_t);

/* Random reading and writing: */
off_t   lseek(int, off_t, int);
ssize_t pread(int, void *, size_t, off_t);
ssize_t pwrite(int, const void *, size_t, off_t);

/* Directories and metadata: */
int     mkdir(const char *, mode_t);
int     rename(const char *, const char *);
int     stat(const char *, struct stat *);
int     unlink(const char *);
```

11 / 15

Run `truss cat <filename>` for effect


Show how C++ `fstream` code calls `fread/fwrite`, which (very indirectly) call `read(2)/write(2)`

File I/O system calls

From a process' perspective:

- system calls are C functions
- files are named by small integers (e.g., FD 3)
- ... but what do these numbers really mean?
- is FD 3 in two processes the same file?

File descriptors



13 / 15

Show some syscall code

What about FD -1?

File descriptors

Small integer numbers

Indices into a kernel array

- the contents of the array are a bit like objects
- like C "objects" in the tutorial!

Each process has *its own* file descriptor array

- my FD 4 != your FD 4